

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
sns.set()
from sklearn.linear_model import LinearRegression
from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split
from sklearn import metrics
from scipy.stats import linregress
from numpy.linalg import inv
```

In [2]:

```
HW1data=pd.read_csv("Q1_data.csv")
```

Question 1a)

In [3]:

```
print(np.mean(HW1data['x1']),np.mean(HW1data['y1']),np.std(HW1data['x1']),np.std(HW1data['y1']),(linregress(HW1data['x1'], HW1data['y1'])))
```

```
9.0 7.500909090909093 3.1622776601683795 1.937024215108669 LinregressResult(slope=0.50009090909090914, intercept=3.0000909090909103, rvalue=0.81642051634483992, pvalue=0.0021696288730787901, stderr=0.11790550059563408)
```

mean of X1: 9.0; mean of Y1: 7.5009; standard deviation of X1: 3.162; standard deviation of y1: 1.937; coefficient of correlation: 0.8164

In [4]:

```
print(np.mean(HW1data['x2']),np.mean(HW1data['y2']),np.std(HW1data['x2']),np.std(HW1data['y2']),(linregress(HW1data['x2'], HW1data['y2'])))
```

```
9.0 7.500909090909091 3.1622776601683795 1.93710869148962 LinregressResult(slope=0.50000000000000011, intercept=3.0009090909090892, rvalue=0.816236506000243, pvalue=0.0021788162369107845, stderr=0.11796374596764074)
```

mean of X2: 9.0; mean of Y2: 7.5009; standard deviation of X2: 3.162; standard deviation of y2: 1.937; coefficient of correlation: 0.8162

In [5]:

```
print(np.mean(HW1data['x3']),np.mean(HW1data['y3']),np.std(HW1data['x3']),np.std(HW1data['y3']),(linregress(HW1data['x3'], HW1data['y3'])))
```

```
9.0 7.5000000000000001 3.1622776601683795 1.9359329439927313 LinregressResult(slope=0.49972727272727291, intercept=3.0024545454545439, rvalue=0.8162867394895984, pvalue=0.0021763052792280152, stderr=0.11787766222100221)
```

mean of X3: 9.0; mean of Y3: 7.5; standard deviation of X3: 3.162; standard deviation of y3: 1.936; coefficient of correlation: 0.8163

In [6]:

```
print(np.mean(HW1data['x4']),np.mean(HW1data['y4']),np.std(HW1data['x4']),np.std(HW1data['y4']),(linregress(HW1data['x4'], HW1data['y4'])))
```

```
9.0 7.5009090909090909 3.1622776601683795 1.9360806451340837 LinregressResult(slope=0.49990909090909091, intercept=3.0017272727272726, rvalue=0.81652143688850276, pvalue=0.0021646023471972222, stderr=0.11781894172968553)
```

mean of X4: 9.0; mean of Y4: 7.5; standard deviation of X4: 3.162; standard deviation of y4: 1.936; coefficient of correlation: 0.8165


Question 1b)

In [7]:

```
X = np.array([[10, 0], [8, 0], [13, 0], [9, 0],[11, 0],[14, 0],[6, 0],[4, 0],[12, 0],[7, 0],[5, 0]])
LinearRegression1 = LinearRegression().fit(X, HW1data['y1'].values)
Slope1=LinearRegression1.coef_[0]
Intercept1=LinearRegression1.intercept_
fig=plt.figure(figsize=(10,5))
plt.scatter(HW1data['x1'], HW1data['y1'], s=100)
plt.ylabel('y1', fontsize = 25)
plt.yticks(fontsize=16)
plt.xlabel('x1', fontsize=25)
plt.xticks(fontsize=16)
Line1 = Slope1*HW1data['x1']+Intercept1
plt.plot(HW1data['x1'],HW1data['y1'],'o',HW1data['x1'],Line1)
```

Out[7]:

```
[<matplotlib.lines.Line2D at 0x1a1bb43be0>,
 <matplotlib.lines.Line2D at 0x1a1bb4a2e8>]
```




In [8]:

```
LinearRegression2 = LinearRegression().fit(X, HW1data['y2'].values)
Slope2=LinearRegression2.coef_[0]
Intercept2=LinearRegression2.intercept_
fig=plt.figure(figsize=(10,5))
plt.scatter(HW1data['x2'], HW1data['y2'], s=100)
plt.ylabel('y2', fontsize = 25)
plt.yticks(fontsize=16)
plt.xlabel('x2', fontsize=25)
plt.xticks(fontsize=16)
Line2 = Slope2*HW1data['x2']+Intercept2
plt.plot(HW1data['x2'],HW1data['y2'],'o',HW1data['x2'],Line2)
```

Out[8]:

```
[<matplotlib.lines.Line2D at 0x1a1bd78978>,
 <matplotlib.lines.Line2D at 0x1a1bd78cc0>]
```




In [9]:

```
LinearRegression3 = LinearRegression().fit(X, HW1data['y3'].values)
Slope3=LinearRegression3.coef_[0]
Intercept3=LinearRegression3.intercept_
fig=plt.figure(figsize=(10,5))
plt.scatter(HW1data['x3'], HW1data['y3'], s=100)
plt.ylabel('y3', fontsize = 25)
plt.yticks(fontsize=16)
plt.xlabel('x3', fontsize=25)
plt.xticks(fontsize=16)
Line3 = Slope3*HW1data['x3']+Intercept3
plt.plot(HW1data['x3'],HW1data['y3'],'o',HW1data['x3'],Line3)
```

Out[9]:

```
[<matplotlib.lines.Line2D at 0x1a1beb66a0>,
 <matplotlib.lines.Line2D at 0x1a1beb69e8>]
```




In [10]:

```
X4=np.array([[8, 0], [8, 0], [8, 0], [8, 0],[8, 0],[8, 0],[8, 0],[19, 0],[8, 0],
[8, 0],[8, 0]])
LinearRegression4 = LinearRegression().fit(X4, HW1data['y4'].values)
Slope4=LinearRegression4.coef_[0]
Intercept4=LinearRegression3.intercept_
fig=plt.figure(figsize=(10,5))
plt.scatter(HW1data['x4'], HW1data['y4'], s=100)
plt.ylabel('y4', fontsize = 25)
plt.yticks(fontsize=16)
plt.xlabel('x4', fontsize=25)
plt.xticks(fontsize=16)
Line4 = Slope4*HW1data['x4']+Intercept4
plt.plot(HW1data['x4'],HW1data['y4'],'o',HW1data['x4'],Line4)
```

Out[10]:

```
[<matplotlib.lines.Line2D at 0x1a1bff0e10>,
 <matplotlib.lines.Line2D at 0x1a1bff0470>]
```



question 1c)

In [12]:

```
y1 = []
y1.append(HW1data['y1'])
y1_pred = LinearRegression1.predict(X)
print(Slope1)
print(Intercept1)
print(np.sum(np.square(np.subtract(y1,np.mean(HW1data['y1'])))))
print(np.sum(np.square(np.subtract(y1,y1_pred))))
print(np.sum(np.square(np.subtract(y1_pred,np.mean(HW1data['y1'])))))
print(np.sum(np.square(np.subtract(y1_pred,np.mean(HW1data['y1'])))/np.sum(np.s
quare(np.subtract(y1,np.mean(HW1data['y1'])))))
```

0.500090909091

3.00009090909

41.2726909091

13.76269

27.5100009091

0.666542459509

For (x1,y1): slope = 0.50009, intercept = 3.00009, SSTO = 41.2727, SSR = 13.763, SSE = 27.51, coefficient of multiple determination = 0.6665

In [13]:

```
y2 = []
y2.append(HW1data['y2'])
y2_pred = LinearRegression2.predict(X)
print(Slope2)
print(Intercept2)
print(np.sum(np.square(np.subtract(y2,np.mean(HW1data['y2'])))))
print(np.sum(np.square(np.subtract(y2,y2_pred))))
print(np.sum(np.square(np.subtract(y2_pred,np.mean(HW1data['y2'])))))
print(np.sum(np.square(np.subtract(y2_pred,np.mean(HW1data['y2'])))/np.sum(np.s
quare(np.subtract(y2,np.mean(HW1data['y2'])))))
```

0.5

3.00090909091

41.2762909091

13.7762909091

27.5

0.666242033727

For (x2,y2): slope = 0.5, intercept = 3.0009, SSTO = 41.2763, SSR = 13.776, SSE = 27.5, coefficient of multiple determination = 0.6662

In [14]:

```
y3 = []
y3.append(HW1data['y3'])
y3_pred = LinearRegression3.predict(X)
print(Slope3)
print(Intercept3)
print(np.sum(np.square(np.subtract(y3,np.mean(HW1data['y3'])))))
print(np.sum(np.square(np.subtract(y3,y3_pred))))
print(np.sum(np.square(np.subtract(y3_pred,np.mean(HW1data['y3'])))))
print(np.sum(np.square(np.subtract(y3_pred,np.mean(HW1data['y3'])))/np.sum(np.s
quare(np.subtract(y3,np.mean(HW1data['y3'])))))
```

0.499727272727

3.00245454545

41.2262

13.7561918182

27.4700081818

0.666324041067

For (x3,y3): slope = 0.4997, intercept = 3.0025, SSTO = 41.2262, SSR = 13.756, SSE = 27.47, coefficient of multiple determination = 0.6663

In [15]:

```
y4 = []
y4.append(HW1data['y4'])
y4_pred = LinearRegression4.predict(X4)
print(Slope4)
print(Intercept4)
print(np.sum(np.square(np.subtract(y4,np.mean(HW1data['y4'])))))
print(np.sum(np.square(np.subtract(y4,y4_pred))))
print(np.sum(np.square(np.subtract(y4_pred,np.mean(HW1data['y4'])))))
print(np.sum(np.square(np.subtract(y4_pred,np.mean(HW1data['y4'])))/np.sum(np.s
quare(np.subtract(y4,np.mean(HW1data['y4'])))))
```

0.499909090909

3.00245454545

41.2324909091

13.74249

27.4900009091

0.666707256898

For (x4,y4): slope = 0.4999, intercept = 3.0025, SSTO = 41.2325, SSR = 13.7425, SSE = 27.49, coefficient of multiple determination = 0.6667

1d) The summary statistics, coefficient of correlation, slope or regression and coefficient of multiple determination are NOT sufficient to judge the quality of fit of linear regression. I made this conclusion because I observed that for all these 4 datasets, all of these values are roughly the same even though the data are very different.

2a)

In [16]:


```
HW2data=pd.read_csv("Q2_data.csv")
```

In [17]:

```
Q2_X1 = np.zeros((500,2))
for i in range(500):
    Q2_X1[i][0] = HW2data['X1'][i]
LinearRegression_2a = LinearRegression().fit(Q2_X1, HW2data['Y'].values)
Slope_2a=LinearRegression_2a.coef_[0]
Intercept_2a=LinearRegression_2a.intercept_
fig=plt.figure(figsize=(10,5))
plt.scatter(HW2data['X1'], HW2data['Y'], s=50)
plt.ylabel('Y', fontsize = 25)
plt.yticks(fontsize=16)
plt.xlabel('X1', fontsize=25)
plt.xticks(fontsize=16)
Line_2a = Slope_2a*HW2data['X1']+Intercept_2a
plt.plot(HW2data['X1'],HW2data['Y'],'o',HW2data['X1'],Line_2a)
```

Out[17]:

```
[<matplotlib.lines.Line2D at 0x1a278a5a20>,  
<matplotlib.lines.Line2D at 0x1a278a5da0>]
```



In [18]:

```
y_2a = []  
y_2a.append(HW2data['Y'])  
y_2a_pred = LinearRegression_2a.predict(Q2_X1)  
print(Slope_2a)  
print(Intercept_2a)  
print(np.sum(np.square(np.subtract(y_2a,np.mean(HW2data['Y'])))))  
print(np.sum(np.square(np.subtract(y_2a,y_2a_pred))))  
print(np.sum(np.square(np.subtract(y_2a_pred,np.mean(HW2data['Y'])))))  
print(np.sum(np.square(np.subtract(y_2a_pred,np.mean(HW2data['Y'])))/np.sum(np.  
square(np.subtract(y_2a,np.mean(HW2data['Y'])))))
```

```
-0.732087336713  
12.1301609462  
4667.38580111  
2151.42326294  
2515.96253817  
0.539051761603
```

The linear model: linear regression. Slope = -0.7321, intercept = 12.1302, SSTO = 4667.39, SSR = 2151.42, SSE = 2515.96, coefficient of multiple determination = 0.53905. The model does NOT fit well.

2b)


In [19]:

```
dataSet1x = []
dataSet1y = []
dataSet2x = []
dataSet2y = []
dataSet3x = []
dataSet3y = []
dataSet4x = []
dataSet4y = []
dataSet5x = []
dataSet5y = []
for i in range(500):
    if HW2data['X2'][i] == 1:
        dataSet1x.append(HW2data['X1'][i])
for i in range(100):
    dataSet1y.append(HW2data['Y'][i])
for i in range(500):
    if HW2data['X2'][i] == 2:
        dataSet2x.append(HW2data['X1'][i])
for i in range(100,200):
    dataSet2y.append(HW2data['Y'][i])
for i in range(500):
    if HW2data['X2'][i] == 3:
        dataSet3x.append(HW2data['X1'][i])
for i in range(200,300):
    dataSet3y.append(HW2data['Y'][i])
for i in range(500):
    if HW2data['X2'][i] == 4:
        dataSet4x.append(HW2data['X1'][i])
for i in range(300,400):
    dataSet4y.append(HW2data['Y'][i])
for i in range(500):
    if HW2data['X2'][i] == 5:
        dataSet5x.append(HW2data['X1'][i])
for i in range(400,500):
    dataSet5y.append(HW2data['Y'][i])
fig=plt.figure(figsize=(20,10))
plt.scatter(dataSet1x,dataSet1y, s=50, color='blue')
plt.scatter(dataSet2x,dataSet2y, s=50, color='yellow')
plt.scatter(dataSet3x,dataSet3y, s=50, color='red')
plt.scatter(dataSet4x,dataSet4y, s=50, color='purple')
plt.scatter(dataSet5x,dataSet5y, s=50, color='green')
plt.ylabel('Y', fontsize = 25)
plt.yticks(fontsize=16)
plt.xlabel('X1', fontsize=25)
plt.xticks(fontsize=16)
```

Out[19]:

```
(array([ -2.,   0.,   2.,   4.,   6.,   8.,  10.,  12.,  14.,  16.]
```

```
',  
<a list of 10 Text xticklabel objects>)
```



2c)

In [20]:

```
dataSet_2c = np.zeros((500,2))
```

In [21]:

```
for i in range(500):  
    dataSet_2c[i][0] = HW2data['X1'][i]
```

In [22]:

```
for i in range(500):  
    dataSet_2c[i][1] = HW2data['X2'][i]
```

In [25]:

```
LinearRegression_2c = LinearRegression().fit(dataSet_2c, HW2data['Y'].values)
```

In [26]:

```
Intercept_2c=LinearRegression_2c.intercept_  
Slope_2c=LinearRegression_2c.coef_
```

In [117]:

```
dataSet1x = []
```

```


dataSet1y = []

dataSet2x = []
dataSet2y = []
dataSet3x = []
dataSet3y = []
dataSet4x = []
dataSet4y = []
dataSet5x = []
dataSet5y = []
for i in range(500):
    if HW2data['X2'][i] == 1:
        dataSet1x.append(HW2data['X1'][i])
for i in range(100):
    dataSet1y.append(HW2data['Y'][i])
for i in range(500):
    if HW2data['X2'][i] == 2:
        dataSet2x.append(HW2data['X1'][i])
for i in range(100,200):
    dataSet2y.append(HW2data['Y'][i])
for i in range(500):
    if HW2data['X2'][i] == 3:
        dataSet3x.append(HW2data['X1'][i])
for i in range(200,300):
    dataSet3y.append(HW2data['Y'][i])
for i in range(500):
    if HW2data['X2'][i] == 4:
        dataSet4x.append(HW2data['X1'][i])
for i in range(300,400):
    dataSet4y.append(HW2data['Y'][i])
for i in range(500):
    if HW2data['X2'][i] == 5:
        dataSet5x.append(HW2data['X1'][i])
for i in range(400,500):
    dataSet5y.append(HW2data['Y'][i])
fig=plt.figure(figsize=(20,10))
plt.ylabel('Y', fontsize = 25)
plt.yticks(fontsize=16)
plt.xlabel('X1', fontsize=25)
plt.xticks(fontsize=16)
Intercept_2c=LinearRegression_2c.intercept_
plt.ylabel('Y', fontsize = 25)
plt.yticks(fontsize=16)
plt.xlabel('X1', fontsize=25)
plt.xticks(fontsize=16)
Line_2c = np.dot(LinearRegression_2c.coef_,dataSet_2c.T)+Intercept_2c
plt.scatter(dataSet1x,dataSet1y, s=50, color='blue')
plt.scatter(dataSet2x,dataSet2y, s=50, color='yellow')
plt.scatter(dataSet3x,dataSet3y, s=50, color='red')
plt.scatter(dataSet4x,dataSet4y, s=50, color='purple')
plt.scatter(dataSet5x,dataSet5y, s=50, color='green')
plt.plot(HW2data['X1'],Line_2c)

```

Out[117]:

[<matplotlib.lines.Line2D at 0x1a29f3a470>]



In [118]:

```
y_2c = []
y_2c.append(HW2data['Y'])
y_2c_pred = LinearRegression_2c.predict(dataSet_2c)
print(Slope_2c)
print(Intercept_2c)
print(np.sum(np.square(np.subtract(y_2c,np.mean(HW2data['Y'])))))
print(np.sum(np.square(np.subtract(y_2c,y_2c_pred))))
print(np.sum(np.square(np.subtract(y_2c_pred,np.mean(HW2data['Y'])))))
print(np.sum(np.square(np.subtract(y_2c_pred,np.mean(HW2data['Y']))))/np.sum(np.
square(np.subtract(y_2c,np.mean(HW2data['Y'])))))
```

```
[ 0.78441638  3.55191332]
-9.15645044033
4667.38580111
331.396183473
4335.98961764
0.928997473619
```

The linear model: multivariate linear regression. Slope = [0.78441638 3.55191332], intercept = -9.1565, SSTO = 4667.4, SSR = 331.4, SSE = 4335.99, coefficient of multiple determination = 0.928997. The model DOES fit well.

2d) The Simpson's paradox: The nature of association between two variables changes when conditioned on another variable. This is true because when we introduced a new predictor X2, we can see that the model fitted five lines instead of one.

Question 3

In [11]:

```
import operator
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import PolynomialFeatures
```

In [59]:

```
x=np.random.uniform(low=-2, high=2, size=(50,))
y=2+3*x+np.random.normal(0,2,50)

x = x[:, np.newaxis]
y = y[:, np.newaxis]


reg_1 = LinearRegression().fit(x,y)

y_pred = reg_1.predict(x)

coef_1 = reg_1.score(x,y)
print(coef_1)

plt.ylabel('Y')
plt.xlabel('X')
plt.scatter(x, y, s=15)
plt.plot(x, y_pred, color='r')
plt.show()
print()
```

0.796351224925



In [60]:

```
x=np.random.uniform(low=-2, high=2, size=(50,))
y=2+3*x+np.random.normal(0,2,50)

x = x[:, np.newaxis]
y = y[:, np.newaxis]


x_degree_2 = PolynomialFeatures(degree=2).fit_transform(x)

reg_2 = LinearRegression().fit(x_degree_2,y)
y_pred_degree_2 = reg_2.predict(x_degree_2)

coef_2 = reg_2.score(x_degree_2,y)
print(coef_2)

plt.scatter(x, y, s=15)
plt.ylabel('Y')
plt.xlabel('X')
sort_axis = operator.itemgetter(0)
sorted_zip = sorted(zip(x,y_pred_degree_2), key=sort_axis)
x, y_pred_degree_2 = zip(*sorted_zip)
plt.plot(x, y_pred_degree_2, color='green')
plt.show()
```

0.771719275592



In [61]:

```
x=np.random.uniform(low=-2, high=2, size=(50,))
y=2+3*x+np.random.normal(0,2,50)

x = x[:, np.newaxis]
y = y[:, np.newaxis]

x_degree_3 = PolynomialFeatures(degree=3).fit_transform(x)


reg_3 = LinearRegression().fit(x_degree_3,y)

y_pred_degree_3 = reg_3.predict(x_degree_3)

coef_3 = reg_3.score(x_degree_3,y)
print(coef_3)

plt.scatter(x, y, s=15)
plt.ylabel('Y')
plt.xlabel('X')
sort_axis = operator.itemgetter(0)
sorted_zip = sorted(zip(x,y_pred_degree_3), key=sort_axis)
x, y_pred_degree_3 = zip(*sorted_zip)
plt.plot(x, y_pred_degree_3, color='blue')
plt.show()
```

0.686463602766



In [62]:

```
x=np.random.uniform(low=-2, high=2, size=(50,))
y=2+3*x+np.random.normal(0,2,50)

x = x[:, np.newaxis]
y = y[:, np.newaxis]

x_degree_4 = PolynomialFeatures(degree=4).fit_transform(x)


reg_4 = LinearRegression().fit(x_degree_4,y)

y_pred_degree_4 = reg_4.predict(x_degree_4)

coef_4 = reg_4.coef_
print(coef_4)

plt.scatter(x, y, s=15)
plt.ylabel('Y')
plt.xlabel('X')
sort_axis = operator.itemgetter(0)
sorted_zip = sorted(zip(x,y_pred_degree_4), key=sort_axis)
x, y_pred_degree_4 = zip(*sorted_zip)
plt.plot(x, y_pred_degree_4, color='yellow')
plt.show()
```

0.775644277324



In [63]:

```
x=np.random.uniform(low=-2, high=2, size=(50,))
y=2+3*x+np.random.normal(0,2,50)

x = x[:, np.newaxis]
y = y[:, np.newaxis]

x_degree_5 = PolynomialFeatures(degree=5).fit_transform(x)


reg_5 = LinearRegression().fit(x_degree_5,y)

y_pred_degree_5 = reg_5.predict(x_degree_5)

coef_5=reg_5.coef_
print(coef_5)

plt.scatter(x, y, s=15)
plt.ylabel('Y')
plt.xlabel('X')
sort_axis = operator.itemgetter(0)
sorted_zip = sorted(zip(x,y_pred_degree_5), key=sort_axis)
x, y_pred_degree_5 = zip(*sorted_zip)
plt.plot(x, y_pred_degree_5, color='orange')
plt.show()
Slope=LinearRegression().fit(x_degree_5,y).coef_
print(Slope)
```

0.785880116742



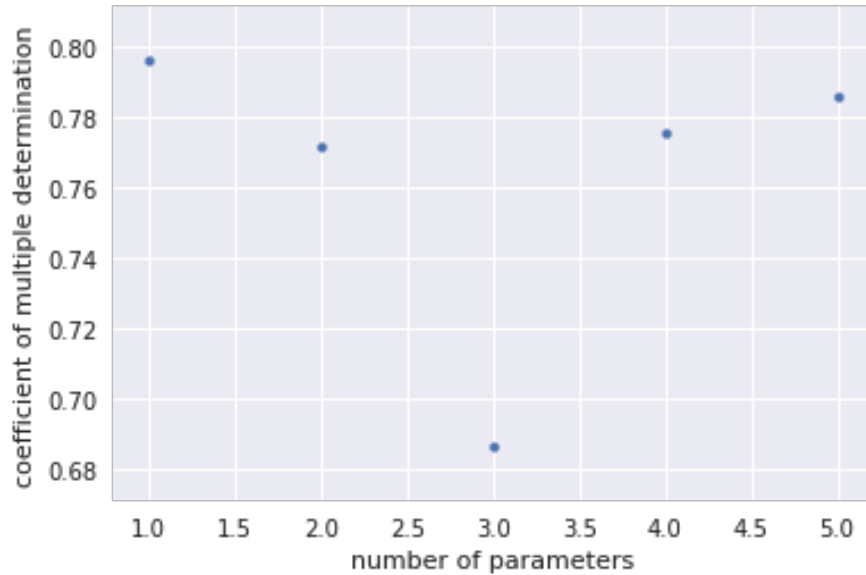
```
[[ 0.          4.21250153 -0.94042119 -0.79545496  0.3117489   0.100
 75729]]
```

In [64]:

```
plt.scatter(x=[1,2,3,4,5], y=[coef_1,coef_2,coef_3,coef_4,coef_5], s=15)
plt.ylabel('coefficient of multiple determination')
plt.xlabel('number of parameters')
```

Out[64]:

Text(0.5, 0, 'number of parameters')



There is no clear correlation between the two.

question 3b)

In []:


```
coef_x = []
slopes=[]
predictions_y=[]
x_vals = [1,1.5,2.25,3.375,5.0625,7.59375]
for w in range(1000):
    x=np.random.uniform(low=-2, high=2, size=(50,))
    y=2+3*x+np.random.normal(0,2,50)
    x = x[:, np.newaxis]
    y = y[:, np.newaxis]
    x_degree_5 = PolynomialFeatures(degree=5).fit_transform(x)
    y_pred_degree_5 = LinearRegression().fit(x_degree_5,y).predict(x_degree_5)
    Slope=LinearRegression().fit(x_degree_5,y).coef_
    slopes.append(Slope)

    coef_x.append(slopes[w][0][1])

    predictions_y.append(np.dot(slopes[-1],x_vals))

plt.hist(coef_x)
plt.xlabel('coefficient associated with X')
plt.ylabel('Frequency')
plt.axvline(x=3.0, color='red')
plt.show()

plt.hist(predictions_y)
plt.xlabel('predictions of Y with X=1.5')
plt.ylabel('Frequency')
plt.axvline(x=6.5, color='red')
plt.show()
```



KeyboardInterrupt
l last)

Traceback (most recent call

<ipython-input-12-8ee5e963a282> in <module>()

```

23 plt.show()
24
---> 25 plt.hist(predictions_y)
26 plt.xlabel('predictions of Y with X=1.5')
27 plt.ylabel('Frequency')

/anaconda3/lib/python3.6/site-packages/matplotlib/pyplot.py in hist(
x, bins, range, density, weights, cumulative, bottom, histtype, align,
orientation, rwidth, log, color, label, stacked, normed, data, **
kwargs)
2617         align=align, orientation=orientation, rwidth=rwidth,
log=log,
2618         color=color, label=label, stacked=stacked, normed=normed,
data=data, **kwargs)
2620
2621 # Autogenerated by boilerplate.py. Do not edit as changes will be lost.

/anaconda3/lib/python3.6/site-packages/matplotlib/__init__.py in inner
er(ax, data, *args, **kwargs)
1783         "the Matplotlib list!") % (label_name, func.__name__),
1784         RuntimeWarning, stacklevel=2)
-> 1785         return func(ax, *args, **kwargs)
1786
1787         inner.__doc__ = _add_data_doc(inner.__doc__,

/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py in hist(self, x, bins, range, density, weights, cumulative, bottom, histtype, align, orientation, rwidth, log, color, label, stacked, normed, **kwargs)
6643         patch = _barfunc(bins[:-1]+boffset, height, width,
6644                          align='center', log=log, color=c, **{bottom_kvar: bottom})
-> 6645         patches.append(patch)
6646         if stacked:
6647             ymin = max(ymin * 0.9, 1e-100)
6648             self.dataLim.intervalx = (ymin, ymax)

/anaconda3/lib/python3.6/site-packages/matplotlib/__init__.py in inner
er(ax, data, *args, **kwargs)
1783         "the Matplotlib list!") % (label_name, func.__name__),
1784         RuntimeWarning, stacklevel=2)
-> 1785         return func(ax, *args, **kwargs)
1786
1787         inner.__doc__ = _add_data_doc(inner.__doc__,

/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py in bar(self, x, height, width, bottom, align, **kwargs)
2332         ymin = max(ymin * 0.9, 1e-100)
2333         self.dataLim.intervaly = (ymin, ymax)

```

```

-> 2334         self.autoscale_view()
    2335
    2336         bar_container = BarContainer(patches, errorbar,
label=label)

/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_base.py in a
utoscale_view(self, tight, scalex, scaley)
    2402             stickies = [artist.sticky_edges for artist in
self.get_children()]
    2403             x_stickies = sum([sticky.x for sticky in
stickies], [])
-> 2404             y_stickies = sum([sticky.y for sticky in
stickies], [])
    2405             if self.get_xscale().lower() == 'log':
    2406                 x_stickies = [xs for xs in x_stickies if xs
> 0]

```

KeyboardInterrupt •

Both of these two graphs above have big variance. They are also slightly biased to the left of the expected values (this means there is an underestimation for both of them). This means that polynomial of degree 5 is not a very good fit. The coefficient of multiple determination, however, told us it is a good fit (0.785880116742). This is why it's misleading.

HOMEWORK 2

4) Ridge regression shrinks the regression coefficients by imposing a penalty on their size. The ridge coefficients minimize a penalized residual sum of squares

$$\hat{\theta}_{\text{Ridge}} = \arg \min_{\theta} \left\{ \sum_{i=1}^N (y_i - \theta_0 - \sum_{j=1}^p x_{ij} \theta_j)^2 + \lambda \sum_{j=1}^p \theta_j^2 \right\}$$

$\lambda \geq 0$ is a complexity parameter that controls the amount of shrinkage.

$$\begin{aligned} \hat{\theta}_{\text{Ridge}} &= \arg \min_{\theta} \left\{ \sum_{i=1}^N (y_i - x_i^T \theta)^2 \right\} + \lambda \|\theta\|_2^2 \\ &= \|Y - X\theta\|_2^2 + \lambda \|\theta\|_2^2 \end{aligned}$$

$$\text{So, } \text{RSS}(\lambda) = (Y - X\theta)^T (Y - X\theta) + \lambda \theta^T \theta$$

$$= (Y^T - \theta^T X^T) (Y - X\theta) + \lambda \theta^T \theta = Y^T Y - Y^T X \theta - \theta^T X^T Y + \theta^T X^T X \theta + \lambda \theta^T \theta$$

$$\begin{aligned} \text{So, } \nabla_{\theta} \text{RSS}(\lambda) &= -(Y^T X)^T - X^T Y + 2X^T X \theta + 2\lambda \theta \\ &= -2X^T Y + 2X^T X \theta + 2\lambda \theta \stackrel{\text{set to}}{=} 0 \end{aligned}$$

$$\text{So, } 2X^T X \theta + 2\lambda \theta = 2X^T Y$$

$$X^T X \theta + \lambda \theta = X^T Y$$

$$\theta (X^T X + \lambda I) = X^T Y$$

$$\text{So, } \hat{\theta}_{\text{Ridge}} = (X^T X + \lambda I)^{-1} X^T Y$$

where I is the $p \times p$ identity matrix

$$5) \text{Var}(\hat{\theta}_{LS}) = \sigma^2 (X^T X)^{-1}$$

$$\text{Var}(\hat{\theta}_{Ridge}) = \sigma^2 (X^T X + \lambda I)^{-1} X^T X (X^T X + \lambda I)^{-1}$$

$$\text{Let } (X^T X) \text{Var}(\hat{\theta}_{LS}) = A$$

$$(X^T X) \text{Var}(\hat{\theta}_{Ridge}) = B$$

$$\text{So } A = \sigma^2 (X^T X)(X^T X)^{-1} = \sigma^2 I$$

$$B = \underbrace{\sigma^2 (X^T X)(X^T X + \lambda I)^{-1}}_C \underbrace{(X^T X)(X^T X + \lambda I)^{-1}}_D$$

$$\text{Let } C = (X^T X)(X^T X + \lambda I)^{-1} = (X^T X + \lambda I - \lambda I)(X^T X + \lambda I)^{-1}$$

$$\text{So } C = (X^T X + \lambda I)(X^T X + \lambda I)^{-1} - \lambda I(X^T X + \lambda I)^{-1}$$

$$= I - \lambda I(X^T X + \lambda I)^{-1}$$

Because $C = D$

So we have the same thing for D.

So, $C \leq I$, and $D \leq I$

$$\text{So, } \sigma^2 C D \leq \sigma^2 I$$

$$\text{So } B \leq A$$

$$\text{So } \text{Var}(\hat{\theta}_{Ridge}) \leq \text{Var}(\hat{\theta}_{LS})$$

Next, look for point $\hat{\theta}$ on the surface to minimize $J(\theta)$

At $\hat{\theta}$, $\nabla J(\theta)$ is orthogonal to the surface (i.e., cannot decrease $J(\theta)$ by moving along the surface)

So, $\nabla J(\theta)$ and $\nabla g(\theta)$ are (anti) parallel vectors.

So there exists λ such that $\nabla J(\theta) + \lambda \nabla g(\theta) = 0$

↓
Lagrange multiplier,
 $\lambda \neq 0$ (any sign)

Lagrangian function: $L(\theta, \lambda) = J(\theta) + \lambda g(\theta)$

Stationary point: $\frac{dL(\theta, \lambda)}{d\theta} = \nabla J(\theta) + \lambda \nabla g(\theta) \xrightarrow{\text{Set to } 0}$

and $\frac{dL(\theta, \lambda)}{d\lambda} = g(\theta) \xrightarrow{\text{Set to } 0}$

$$L(\theta, \lambda) = \|y - X\theta\|_2^2 + \lambda (W^T \theta - b)$$

$$\frac{dL(\theta, \lambda)}{d\theta} = 2X^T(X\theta - y) + \lambda W \xrightarrow{\text{Set to } 0} \quad (1)$$

$$\frac{dL(\theta, \lambda)}{d\lambda} = W^T \theta - b \xrightarrow{\text{Set to } 0} \quad (2)$$

from (1), $2X^T X \theta - 2X^T y + \lambda W = 0$

$$2X^T X \theta - 2X^T y = -\lambda W$$

$$2X^T X \theta = 2X^T y - \lambda W$$

$$X^T X \theta = X^T y - \frac{1}{2} \lambda W$$

$$\hat{\theta} = (X^T X)^{-1} (X^T y - \frac{1}{2} \lambda W) = X^T y - \frac{1}{2} \lambda W$$

because $W^T \hat{\theta} = b$

$$\text{So } W^T (X^T Y - \frac{1}{2} \lambda W) = b$$

$$W^T X^T Y - \frac{1}{2} \lambda W^T W = b$$

$$\frac{1}{2} \lambda W^T W = W^T X^T Y - b$$

$$\lambda = 2 (W^T W)^{-1} (W^T X^T Y - b)$$

$$\text{So, } \hat{\theta} = X^T Y - \frac{1}{2} 2 (W^T W)^{-1} (W^T X^T Y - b) W$$

$$= X^T Y - (W^T W)^{-1} (W^T X^T Y - b) W$$

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
sns.set()
from sklearn.linear_model import LinearRegression
from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import train_test_split
from sklearn import metrics
from scipy.stats import linregress
from numpy.linalg import inv
import operator
import matplotlib.pyplot as plt
import matplotlib
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import PolynomialFeatures
```

Question 7a)i)

In []:

```
theta_7ai= np.dot(inv(np.dot(X_7ai.T,X_7ai)),np.dot(X_7ai.T,Y_7ai))
```

Question 7a)ii)

In []:

```
theta_7aai = np.dot(inv(np.dot(X_7aai.T,X_7aai)+regularization_parameter*np.identity(number_j)),np.dot(X_7aai.T,Y_7aai))
```

Question 7a)iii)

In [2]:

```
def gradient_descent(x, y, initial_theta, regularization_parameter, loss_func, alpha=0.01, precision=0.001):
    loss = []
    theta = initial_theta
    all_thetas = [] # to store all thetas
    predictions = [] # to store all predictions
    number_of_steps = 0
    previous_loss = 0
    prediction = np.dot(x, theta) #dot product
    error = prediction - y
    current_loss = loss_func(error, theta, regularization_parameter)
    predictions.append(prediction)
    loss.append(current_loss)
    all_thetas.append(theta)
    number_of_steps += 1
    while abs(current_loss - previous_loss) > precision: #if the difference between current and previous values of loss function is bigger than the precision we set
        previous_loss = current_loss #we update the value of the loss function to be the current value of loss function
        gradient = np.dot(x.T, error / np.size(y)) #new gradient
        theta = theta - alpha * gradient #update theta
        all_thetas.append(theta)

        prediction = np.dot(x, theta)
        error = prediction - y
        current_loss = loss_func(error, theta, regularization_parameter)
        loss.append(current_loss)

    return all_thetas, loss, predictions

def loss_func(error, theta, regularization_parameter): #lasso loss
    return np.sum(error**2) + regularization_parameter * np.sum(np.abs(theta))
```

Question 7b)i) - analytical solution for least squares regression


In [3]:

```
x=np.random.uniform(low=-2, high=2, size=(50,))
y=2+3*x+np.random.normal(0,2,50)

A=np.ones(50)
B=np.array((x))
C=np.array((x**2))
D=np.array((x**3))
E=np.array((x**4))
F=np.array((x**5))
matrix_x = np.column_stack((A,B,C,D,E,F))

theta_7bi= np.dot(inv(np.dot(matrix_x.T,matrix_x)),np.dot(matrix_x.T,y))
y_analytical_solution=np.dot(matrix_x,theta_7bi)

plt.scatter(x, y, s=10)
sort_axis = operator.itemgetter(0)
sorted_zip = sorted(zip(x,y_analytical_solution), key=sort_axis)
x_copy, y_analytical_zip = zip(*sorted_zip)
plt.plot(x_copy, y_analytical_zip)
plt.xlabel('X')
plt.ylabel('Y')
plt.show()
```



Question 7b)ii) - analytical solution for ridge regression as function of the regularization parameter

In [4]:

```
regularization_parameter=0.5
```

```
theta_7bii = np.dot(inv(np.dot(matrix_x.T,matrix_x)+regularization_parameter*np.  
identity(6)),np.dot(matrix_x.T,y))
```

```
y_ridge_regression=np.dot(matrix_x,theta_7bii)
```

```
plt.scatter(x, y, s=10)
```

```
sort_axis = operator.itemgetter(0)
```

```
sorted_zip = sorted(zip(x,y_ridge_regression), key=sort_axis)
```


```
x_copy, y_ridge_zip = zip(*sorted_zip)
```

```
plt.plot(x_copy, y_ridge_zip)
```

```
plt.xlabel('X')
```

```
plt.ylabel('Y')
```

```
plt.show()
```




Question 7b)iii) - batch gradient descent for lasso regression

In [5]:

```
regularization_parameter=0.5

theta_i = np.random.rand(6)
all_thetas, loss, predictions = gradient_descent(matrix_x, y, theta_i, regularization_parameter, loss_func)

y_lasso_regression=np.dot(matrix_x,all_thetas[-1])
plt.scatter(x, y, s=10)
sort_axis = operator.itemgetter(0)
sorted_zip = sorted(zip(x,y_lasso_regression), key=sort_axis)
x_copy, y_lasso_zip = zip(*sorted_zip)
plt.plot(x_copy, y_lasso_zip)
plt.xlabel('X')
plt.ylabel('Y')
plt.show()
```



Question 7c) – analytical solution for ridge regression

In [12]:


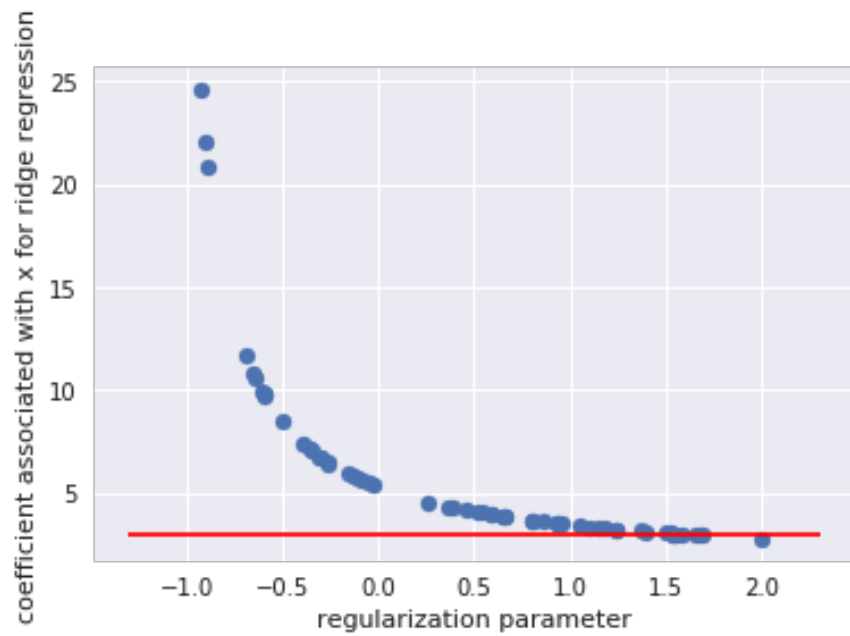
```
regularization_parameter=np.random.uniform(low=-1, high=2, size=(60,))
all_theta_7c_ridge=[]
coef_associated_with_x_ridge=[]
predictions_y_ridge=[]
x_vals = [1,1.5,2.25,3.375,5.0625,7.59375]
for q in range(60):
    theta_7c_ridge = np.dot(inv(np.dot(matrix_x.T,matrix_x)+regularization_parameter[q]*np.identity(6)),np.dot(matrix_x.T,y))
    all_theta_7c_ridge.append(theta_7c_ridge)

    coef_associated_with_x_ridge.append(all_theta_7c_ridge[q][1])

    predictions_y_ridge.append(np.dot(all_theta_7c_ridge[-1],x_vals))

plt.scatter(regularization_parameter,coef_associated_with_x_ridge)
plt.xlabel('regularization parameter')
plt.ylabel('coefficient associated with x for ridge regression')
plt.hlines(y=3,xmin=-1.3,xmax=2.3,color='red')
plt.show()

plt.scatter(regularization_parameter,predictions_y_ridge)
plt.xlabel('regularization parameter')
plt.ylabel('predictions of Y for X=1.5 for ridge regression')
plt.hlines(y=6.5,xmin=-1.3,xmax=2.3,color='red')
plt.show()
```



Question 7c) batch gradient descent optimization for lasso regression


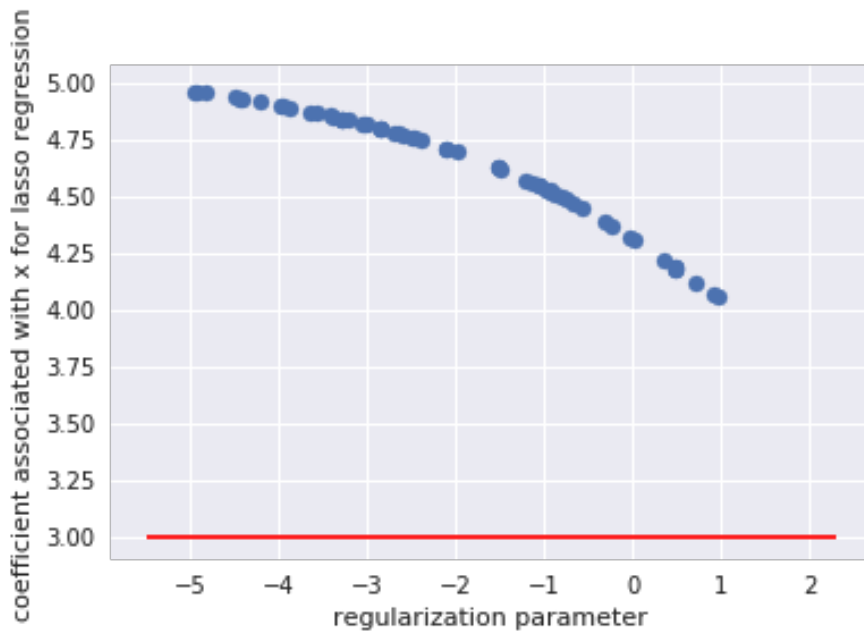
In [13]:

```
regularization_parameter=np.random.uniform(low=-5, high=1, size=(60,))
all_theta_7c_lasso=[]
coef_associated_with_x_lasso=[]
predictions_y_lasso=[]
x_vals = [1,1.5,2.25,3.375,5.0625,7.59375]
for q in range(60):
    theta_i = np.random.rand(6)
    all_thetas, loss, predictions = gradient_descent(matrix_x, y, theta_i, regularization_parameter[q], loss_func)
    coef_associated_with_x_lasso.append(all_thetas[-1][1])

    predictions_y_lasso.append(np.dot(all_thetas[-1],x_vals))

plt.scatter(regularization_parameter,coef_associated_with_x_lasso)
plt.xlabel('regularization parameter')
plt.ylabel('coefficent associated with x for lasso regression')
plt.hlines(y=3,xmin=-5.5,xmax=2.3,color='red')
plt.show()

plt.scatter(regularization_parameter,predictions_y_lasso)
plt.xlabel('regularization parameter')
plt.ylabel('predictions of Y for X=1.5 for lasso regression')
plt.hlines(y=6.5,xmin=-5.5,xmax=2.3,color='red')
plt.show()
```



For analytical solution for ridge regression, parameters are bigger when regularization parameter is negative (overestimation). When regularization parameter gets bigger or positive, parameters get exponentially smaller and asymptotically approach the true value.

For batch gradient descent using Lasso regression, the parameters are overestimated for negative/small positive regularization parameter. As the regularization parameter gets bigger, the parameter and slope approach the true value.

Question 7d)i – analytical solution for least squares

In [15]:



```
coef_with_x_7di = []
all_theta_7d_least_sq=[]
predictions_y_7di=[]
x_vals = [1,1.5,2.25,3.375,5.0625,7.59375]
for w in range(1000):
    x=np.random.uniform(low=-2, high=2, size=(50,))
    y=2+3*x+np.random.normal(0,2,50)
    A=np.ones(50)
    B=np.array((x))
    C=np.array((x**2))
    D=np.array((x**3))
    E=np.array((x**4))
    F=np.array((x**5))
    matrix_x = np.column_stack((A,B,C,D,E,F))
    theta_7di= np.dot(inv(np.dot(matrix_x.T,matrix_x)),np.dot(matrix_x.T,y))
    all_theta_7d_least_sq.append(theta_7di)

    coef_with_x_7di.append(all_theta_7d_least_sq[w][1])

    predictions_y_7di.append(np.dot(all_theta_7d_least_sq[-1],x_vals))

plt.hist(coef_with_x_7di)
plt.xlabel('coefficient associated with X')
plt.ylabel('Frequency')
plt.axvline(x=3.0, color='red')
plt.show()

plt.hist(predictions_y_7di)
plt.xlabel('predictions of Y with X=1.5')
plt.ylabel('Frequency')
plt.axvline(x=6.5, color='red')
plt.show()
```



Unbiased but large variance.

Question 7d)ii – analytical solution for ridge regression

In [17]:



```
coef_with_x_7dii = []
all_theta_7d_ridge=[]
predictions_y_7dii=[]
x_vals = [1,1.5,2.25,3.375,5.0625,7.59375]
for w in range(1000):
    x=np.random.uniform(low=-2, high=2, size=(50,))
    y=2+3*x+np.random.normal(0,2,50)
    A=np.ones(50)
    B=np.array((x))
    C=np.array((x**2))
    D=np.array((x**3))
    E=np.array((x**4))
    F=np.array((x**5))
    matrix_x = np.column_stack((A,B,C,D,E,F))
    theta_7d_ridge = np.dot(inv(np.dot(matrix_x.T,matrix_x)+0.5*np.identity(6)),
np.dot(matrix_x.T,y))
    all_theta_7d_ridge.append(theta_7d_ridge)

    coef_with_x_7dii.append(all_theta_7d_ridge[w][1])

    predictions_y_7dii.append(np.dot(all_theta_7d_ridge[-1],x_vals))

plt.hist(coef_with_x_7dii)
plt.xlabel('coefficient associated with X')
plt.ylabel('Frequency')
plt.axvline(x=3.0, color='red')
plt.show()

plt.hist(predictions_y_7dii)
plt.xlabel('predictions of Y with X=1.5')
plt.ylabel('Frequency')
plt.axvline(x=6.5, color='red')
plt.show()
```



A bit biased to the left (underestimation) and big variance.

Question 7diii bath gradient descent for Lasso

In [18]:


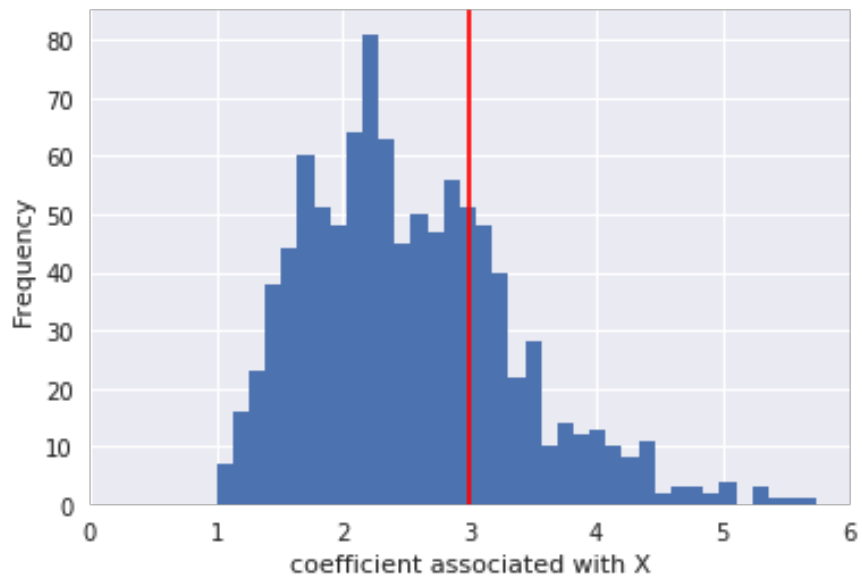
```
coef_with_x_7diii = []
all_theta_7d_lasso=[]
predictions_y_7diii=[]
x_vals = [1,1.5,2.25,3.375,5.0625,7.59375]
for w in range(1000):
    x=np.random.uniform(low=-2, high=2, size=(50,))
    y=2+3*x+np.random.normal(0,2,50)
    A=np.ones(50)
    B=np.array((x))
    C=np.array((x**2))
    D=np.array((x**3))
    E=np.array((x**4))
    F=np.array((x**5))
    matrix_x = np.column_stack((A,B,C,D,E,F))
    theta_i = np.random.rand(6)
    all_thetas, loss, predictions = gradient_descent(matrix_x, y, theta_i, 0.5,
loss_func)
    coef_with_x_7diii.append(all_thetas[-1][1])

    predictions_y_7diii.append(np.dot(all_thetas[-1],x_vals))

plt.hist(coef_with_x_7diii, bins=np.linspace(1,6,40))
plt.xlabel('coefficient associated with X')
plt.ylabel('Frequency')
plt.xlim([0,6])
plt.axvline(x=3.0, color='red')
plt.show()

plt.hist(predictions_y_7diii, bins=np.linspace(3,10,40))
plt.xlabel('predictions of Y with X=1.5')
plt.ylabel('Frequency')
plt.axvline(x=6.5, color='red')
plt.show()
```

```
/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:15: Run
timeWarning: invalid value encountered in double_scalars
from ipykernel import kernelapp as app
```



The graphs show a large variance. It is also biased towards underestimating both the slope of the X term.